

# Real-time Analytics on Medical Device Data

Brock Noland

brock@phdata.io  
@phdatainc  
phdata.io



# About phData

---

Managed Services Provider fully dedicated to Hadoop, Spark, and Kafka



## analytics

---

Tackle your organization's most challenging analytical problems.



## applications

---

Operationalize Hadoop and bring your Hadoop-driven application into production.



## infrastructure

---

Deploy, develop and adopt Hadoop and its supporting technologies.

# About Brock Noland

---

- Co-founder and Chief Architect for phData
- Co-founder of Apache Sentry, Vice President of Apache MRUnit, PMC on Apache Hive, Flume, Crunch, Parquet, Incubator, and Apache Member
- Apache mentor to Apache Kudu and Impala (incubating)
- Early Clouderan and earlier Hadoop user (Thomson Reuters)
- Trainer, Solution Architect, Engineer, and Engineering Manager
- Deep Spark experience
  - Hive on Spark
  - StreamSets

- **Current trends**
- Medical devices
- Infrastructure
- Outcome



# Current Trends

- Mobile technology miniaturization
- Always-on internet connections
- Big data technologies


# Mobile technology miniaturization



# Always-on internet connections

## Mobile Web Usage Growing

**Forward Projection: Mobile Web Browsing vs. Desktop Web Browsing**  
(2007-2015)



Source: Mary Meeker, Morgan Stanley, "Internet Trends," April 12, 2010



# Big data technologies

- Easily ingest, analyze and manage massive data volumes
- Previous generation technologies (RDBMS) capped out in low Terabytes
- Apache Hadoop/Spark installs start in low Terabytes



# Convergence

***"More specifically, Rubin believes Internet-connected devices (smartphones, tablets, thermostats, smoke detectors, and cars, for example) will create massive amounts of data that will be analyzed by deep-learning technologies."***

Andy Rubin: AI Is The Future Of Computing, Mobility - <http://goo.gl/iJCnu8>



- Current trends
- **Medical devices**
- Infrastructure
- Outcome

What does this mean for medical devices?





# Medical devices

- Internal (Implantable)
- External

# Internal medical devices



# External medical devices

- e.g. Cardiac Event Recorders (CER)
- Rechargeable batteries
- Trivial to iterate
- Can be as simple as a watch




- Current trends
- Medical devices
- **Infrastructure**
- Outcome

Devices



Collectors



Kafka



Hadoop



Analytics



# Kafka



# Hadoop

Stream Processing



SQL



# Analytics





# Apache Kafka

- Open source, scalable message system
- Important for real-time IoT use cases
- 1.1 trillion messages per day at LinkedIn  
<http://goo.gl/Wqb72H>

# Apache Hadoop


- Open source, scalable **ecosystem**
- Important for storing large historical data sets
- 455 petabytes at Yahoo  
<http://goo.gl/utc4qM>

# What is the Hadoop ecosystem?

- Apache
  - Atlas
  - Ambari
  - Accumulo
  - Avro
  - Beam
  - Bigtop
  - Chukwa
  - Crunch
  - Eagle
  - Falcon
  - Flink
  - Flume
  - Giraph
  - Hadoop
    - HDFS
    - MapReduce
    - YARN
  - HBase
  - Hama
  - Hive
  - Impala
  - Incubator
- Apache
  - Kylin
  - Kudu
  - Kafka
  - Knox
  - Mesos
  - Myriad
  - Metron
  - Mahout
  - Nifi
  - Oozie
  - ORC
  - Oozie
  - Parquet
  - Pig
  - REEF
  - Ranger
  - Sentry
  - Storm
  - Samza
  - Slider
  - Solr
  - Spark
- Apache
  - Spot
  - Sqoop
  - SystemML
  - Twill
  - Thrift
  - Whir
  - Zeppelin
  - Zookeeper
- AtScale
- Alluxio/Tachyon
- Cask
- DataFu
- Dremel
- Kylin
- Hue
- BigTable
- Deeplearning4j
- H2O
- Oryx
- StreamSets
- Sense



# HDFS



Doug &  
Mike

- Started building a web search engine in 2002
  - Crawl entire internet
  - Generate large files
  - Process and index huge volumes of data
- Discovered Google's GFS paper in 2002, MapReduce in 2004

# HDFS Design Assumptions

- Files are append-only
- Files are large (GBs to TBs)
- Accesses are large and sequential

# HDFS Today

- HDFS is great at:
  - Large scans
  - Unstructured data
  - Batch ingest
  - Ingest/access of large files
- HDFS is NOT great at:
  - Random read/write
  - Real-time ingest
  - Updates



# HBase

---

## **Bigtable: A Distributed Storage System for Structured Data**

Fay Chang, Jeffrey Dean, Sanjay Ghemawat, Wilson C. Hsieh, Deborah A. Wallach  
Mike Burrows, Tushar Chandra, Andrew Fikes, Robert E. Gruber

{fay,jeff,sanjay,wilsonh,kerr,m3b,tushar,fikes,gruber}@google.com

*Google, Inc.*

# HBase Design

- Need OLTP-like storage
  - Low latency
  - Keys are indexed
  - Provide fast random read/write access
  - Mutable
- Same properties apply to Cassandra

# Apache Spark

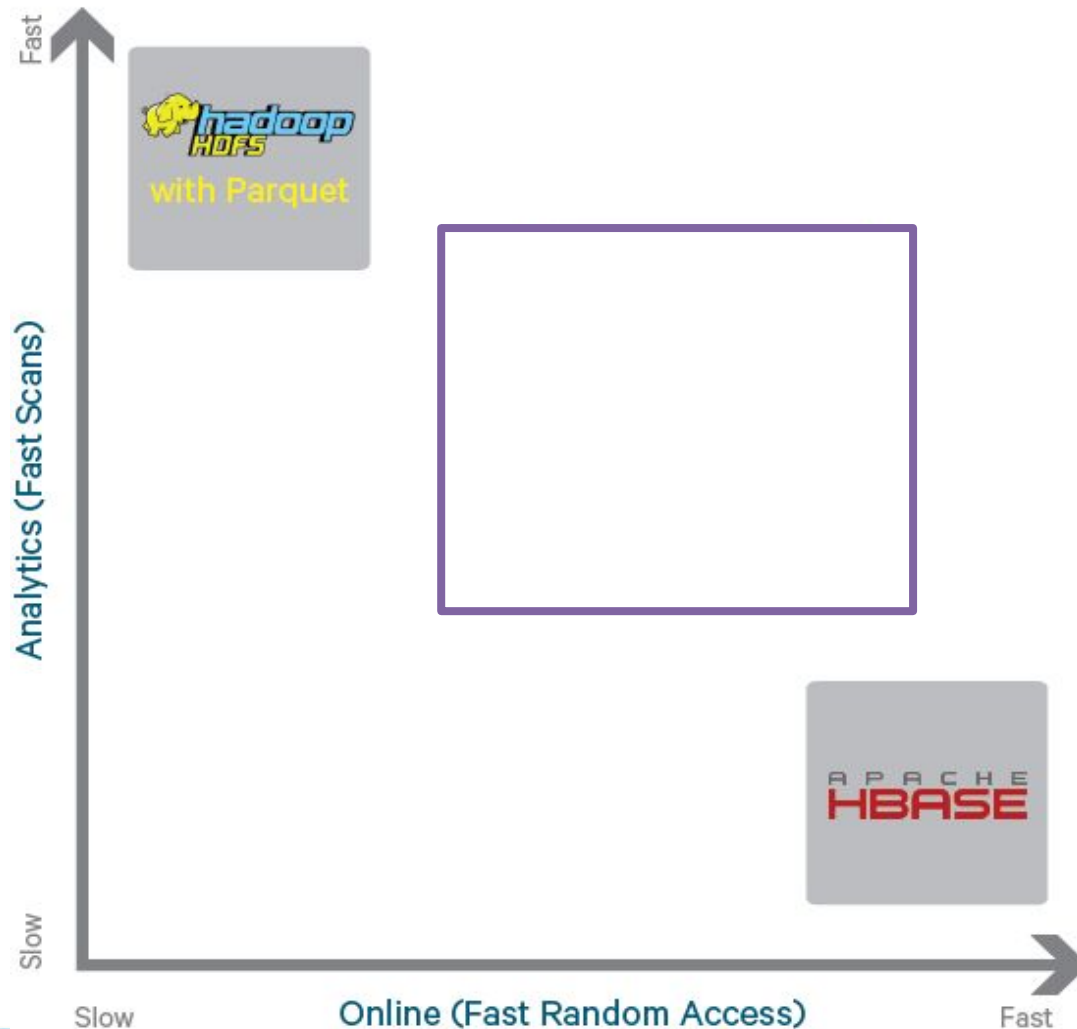
- Open source, scalable processing engine
- Important for scalable stream processing (among other things)
- Processing all daily rides at Uber  
<http://goo.gl/dclAub>



# SQL

- Many projects providing SQL
  - Apache Hive
  - Apache SparkSQL
  - Cloudera Implala
- Important for post hoc analytics

# Previous storage landscape of the Hadoop ecosystem



**HDFS** (GFS) excels at:

- Batch ingest only (eg hourly)
- Efficiently scanning large amounts of data (analytics)

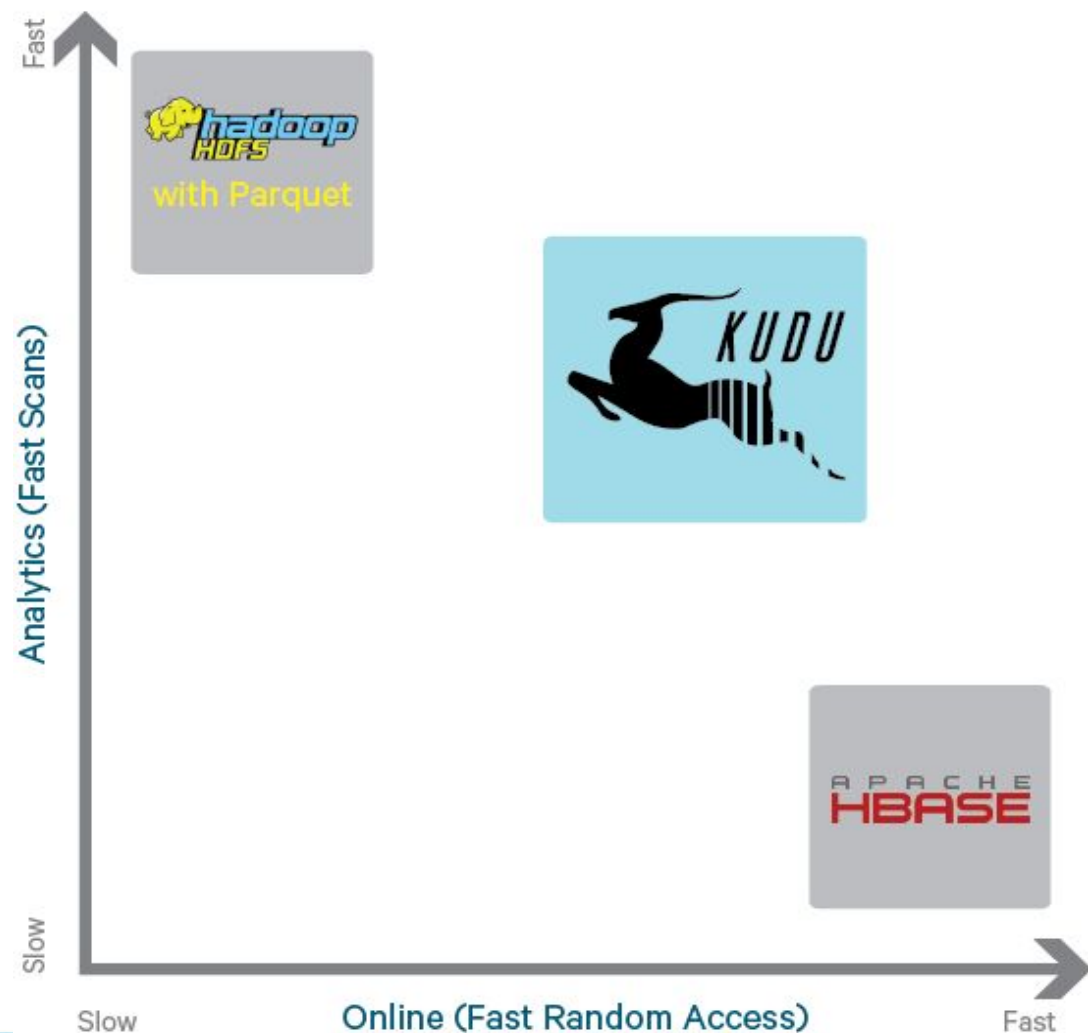
**HBase** (BigTable) excels at:

- Efficiently finding and writing individual rows
- Making data mutable

Gaps exist when these properties are needed *simultaneously*

- HDFS allows for fast writes and scans, but updates are slow and cumbersome
- HBase is fast for updates and inserts at the expense of data scans (i.e. analytics)

# Kudu design goals



- **High throughput** for big scans  
*Goal: Within 2x of Parquet*
- **Low-latency** for short accesses  
*Goal: 1ms read/write on SSD*
- **Database-like** semantics  
(initially single-row ACID)
- **Relational data model**
  - SQL queries are easy
  - “NoSQL” style scan/insert/update (Java/C++ client)



# What Kudu is ***NOT***

Not a SQL interface itself

- ✓ It's a storage layer

Not an application that runs on HDFS

- ✓ It's an alternative, native Hadoop storage engine

Not a replacement for HDFS or HBase

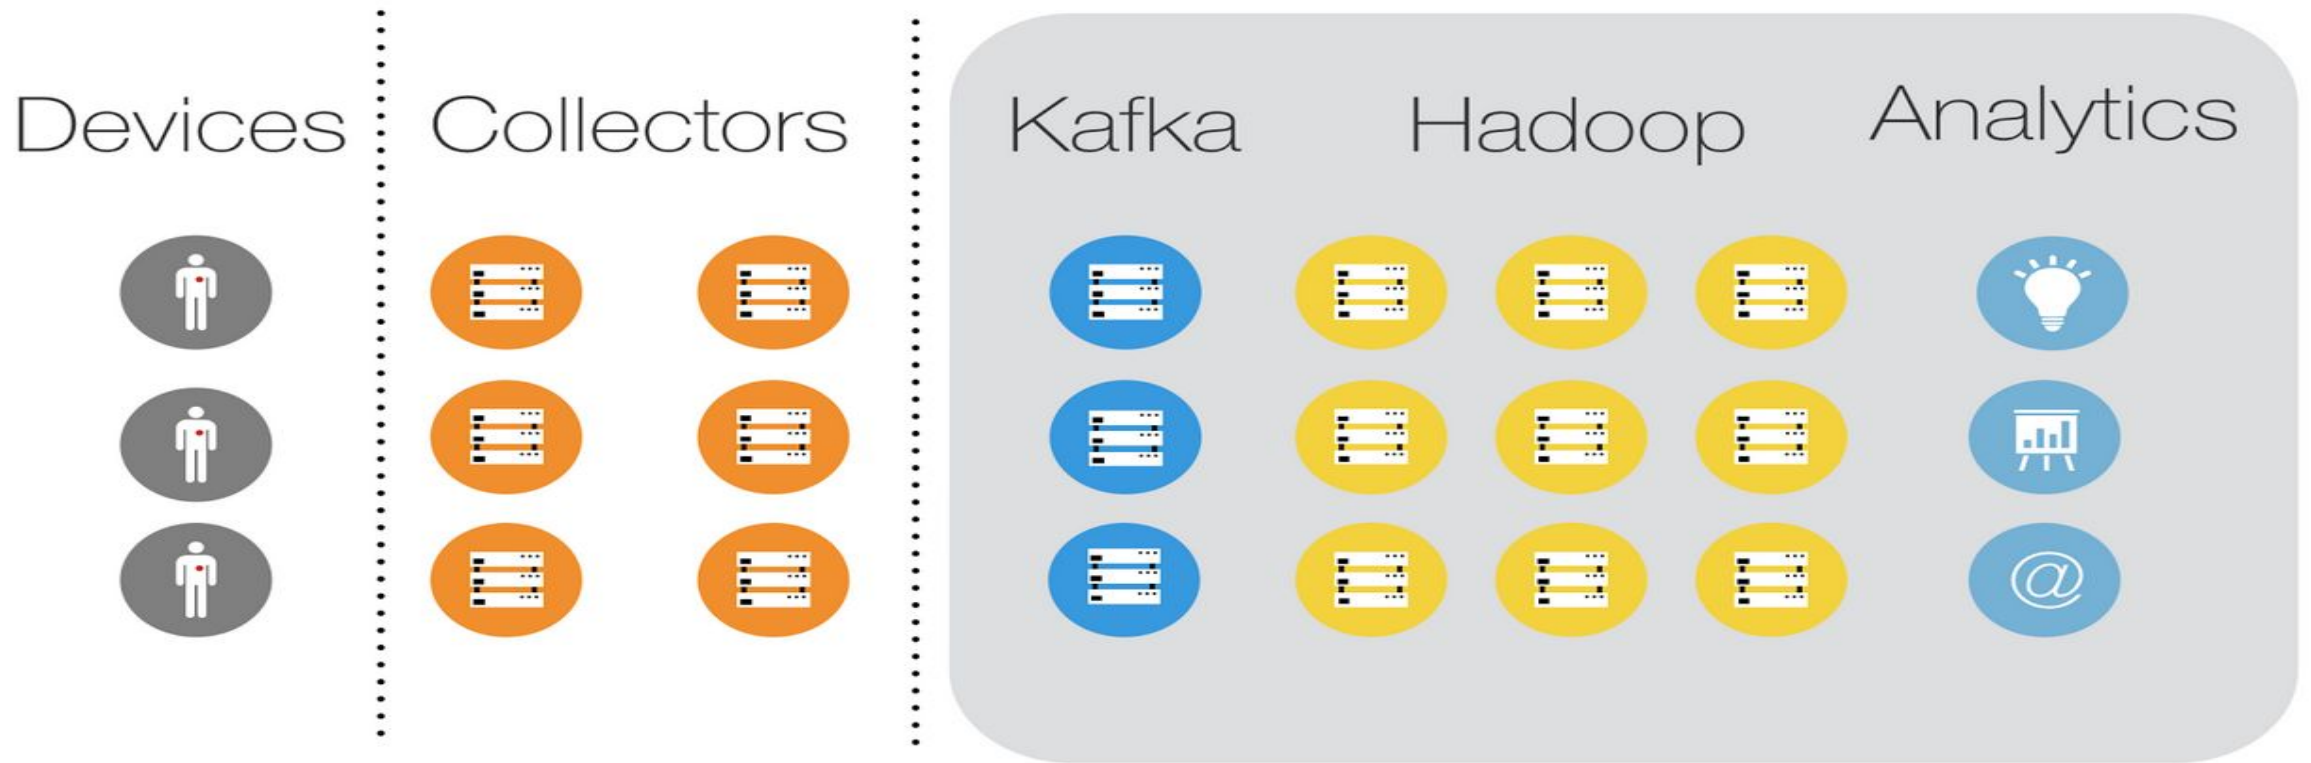
- ✓ Select the right storage for the right use case
- ✓ Cloudera will continue to support and invest in all three

# Common Kudu use cases

**Kudu is best for use cases requiring a simultaneous combination of sequential and random reads and writes**

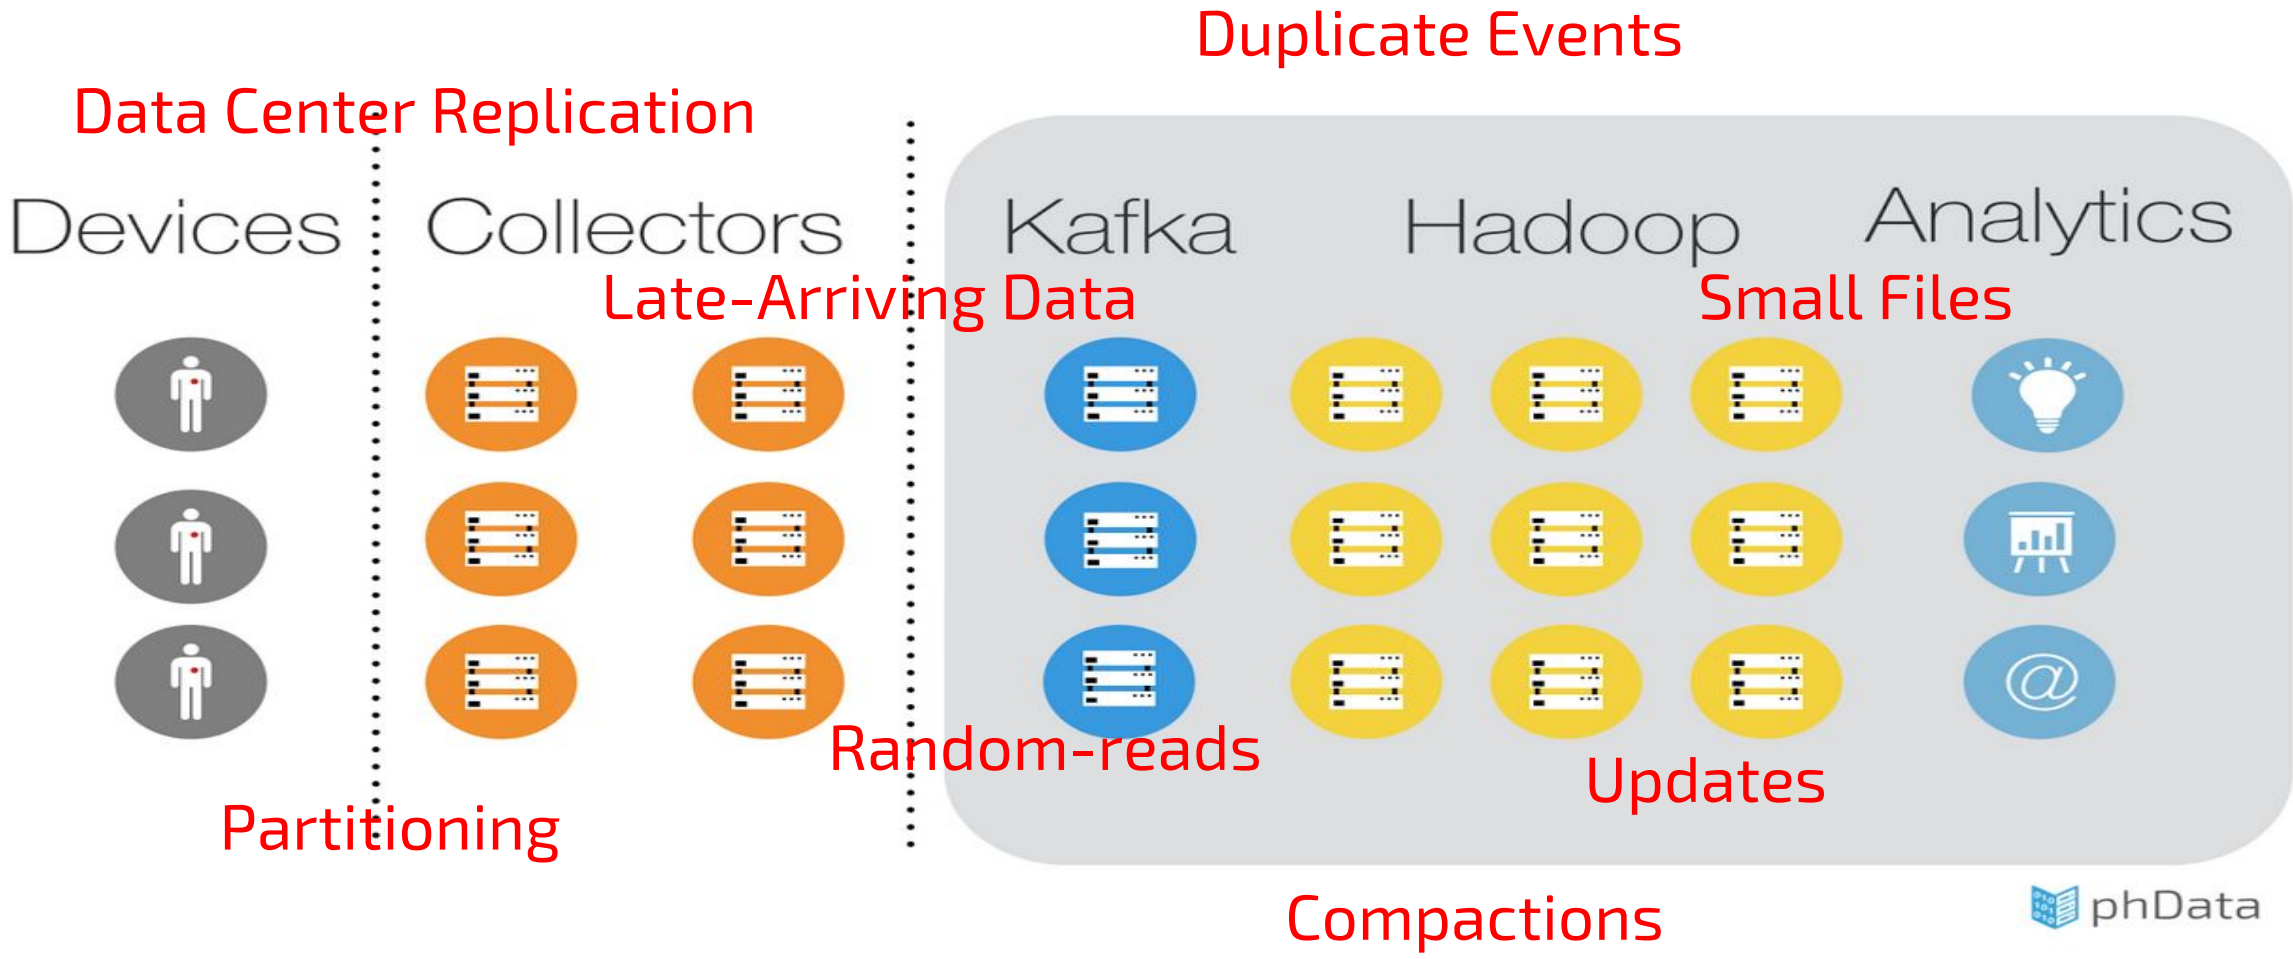
- **Time series**
  - Examples: Streaming market data, fraud detection / prevention, risk monitoring
  - Workload: Insert, updates, scans, lookups
- **Machine data analytics**
  - Example: Network threat detection
  - Workload: Inserts, scans, lookups
- **Online reporting**
  - Example: Operational data store (ODS)
  - Workload: Inserts, updates, scans, lookups
- **Potentially anything where the data is not append only** (or needs scale, parallelism, ...)

# How would we build any Streaming or IoT Analytics System Today?





# What makes this hard?



## Real-time Analytics on Medical Device Data – Part 1 – Introduction

BY: [BROCK NOLAND](#) | JULY 21, 2015 | [HADOOP](#), [INGEST](#), [KAFKA](#), [MEDICAL](#), [SQL](#)

*This is the first post in a series of posts on real-time analytics on medical device data. Read post two on the required [infrastructure](#) and post three on the [SQL schema](#).*

phData believes the Internet of Things (IoT) and Big Data technology such as Apache Hadoop will result in big changes in health care. Dramatic advances in [Telemedicine](#) will be made possible by the constant miniaturization of electronics, ever present internet connections and the ability to store and analyze large volumes of data.

## Raw Data

date	t1	patient_id	a	...	z
20150101	00:00:01:43	12345	1.4	...	0.04
20150101	00:00:01:12	98273	1.7	...	0.12
20150101	00:00:02:43	12345	1.4	...	0.06
20150102	00:00:03:43	12345	1.3	...	0.07
20150102	00:00:01:37	58796	2.1	...	0.07
...	...	...	...	...	...
20151231	11:59:59:43	12345	1.6	...	0.04
20151231	11:59:59:12	98273	1.1	...	0.11



**365\*24\*60\*10 million records per year**

## Patient ID: 12345

date	t1	a	...	z
20150101	00:00:01:43	1.4	...	0.04
20150101	00:00:02:43	1.4	...	0.06
20150102	00:00:03:43	1.3	...	0.07
...	...	...	...	...
20151231	11:59:59:43	1.6	...	0.04

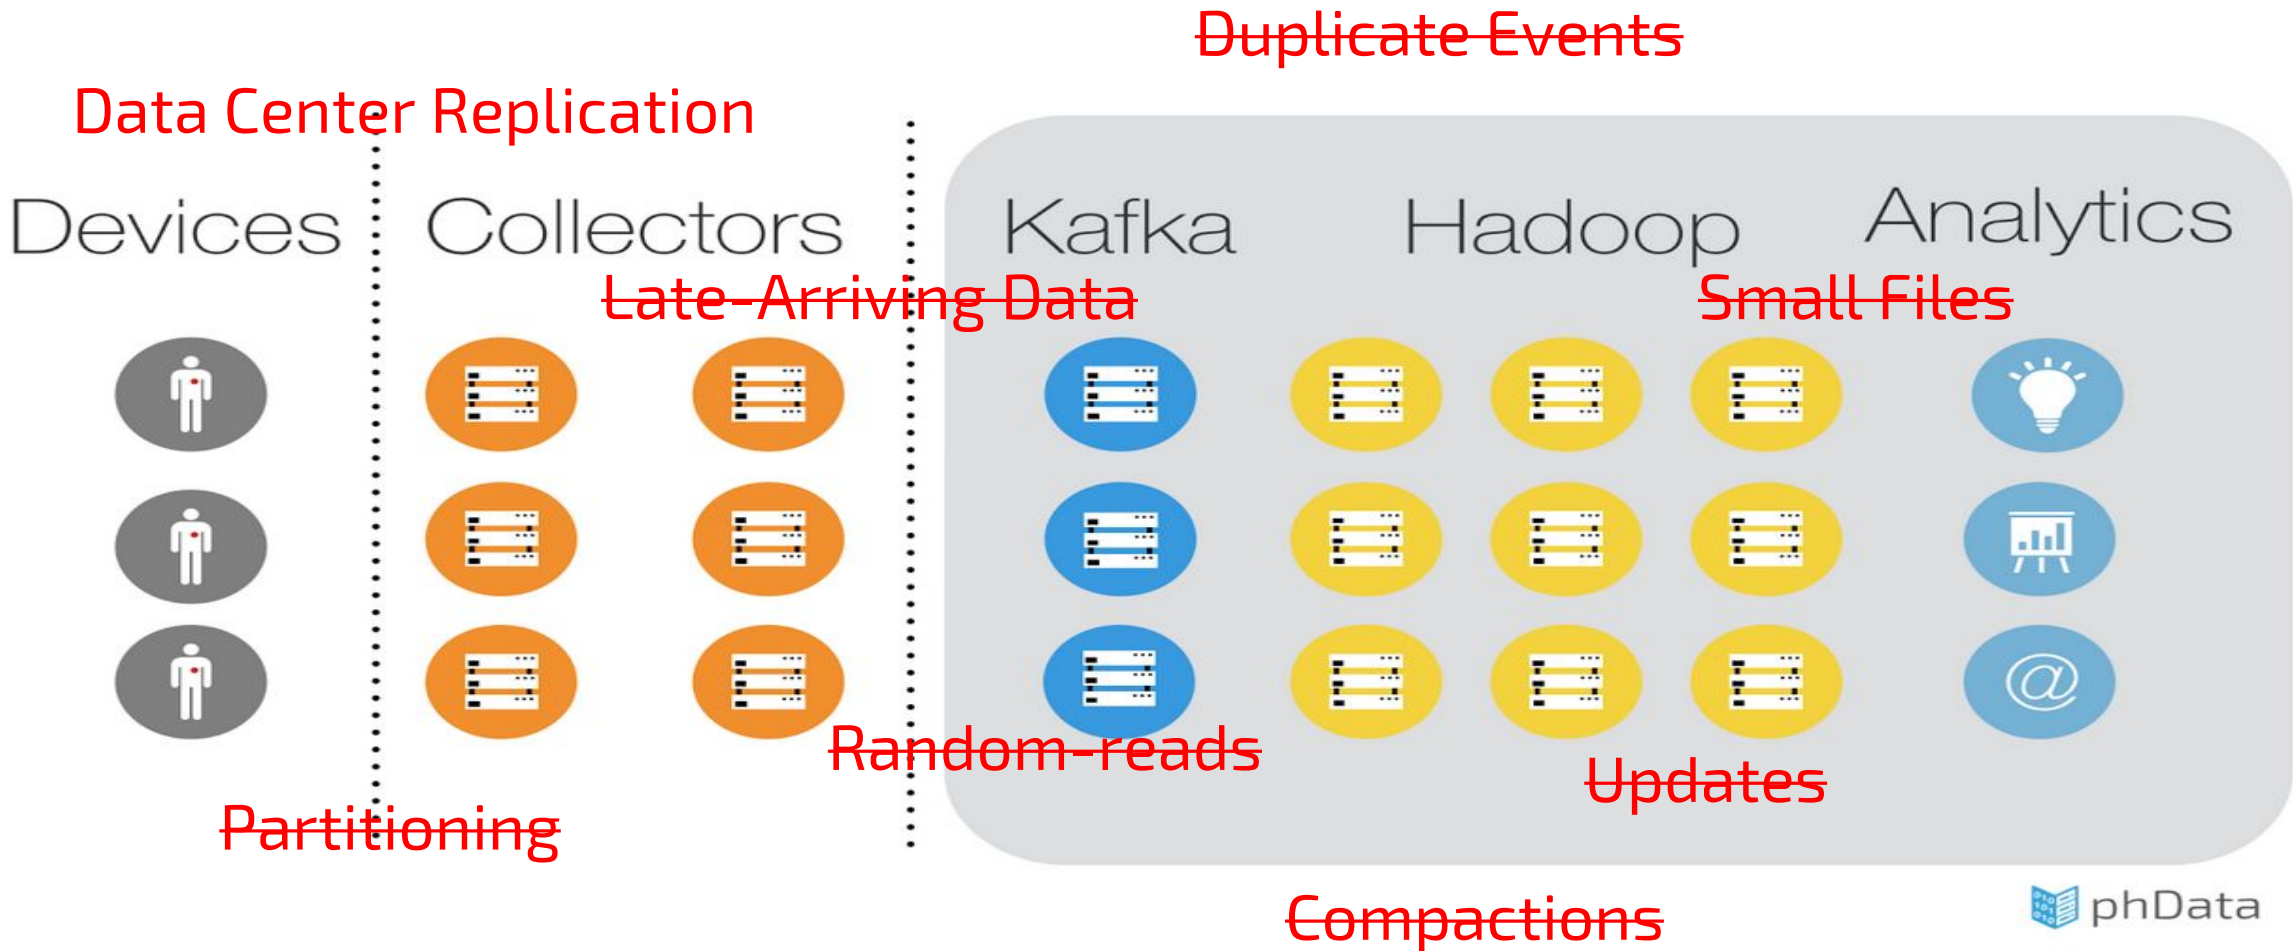


**525,600 records per patient per year**

## Partition by Date

date	t1	patient_id	a	...	z
20150101	00:00:01:43	12345	1.4	...	0.04
20150101	00:00:02:43	12345	1.4	...	0.06
...	...	...	...	...	...
20150101	11:59:59:43	12345	1.3	...	0.04
20150102	00:00:01:43	12345	1.4	...	0.05
20150102	00:00:02:43	12345	1.5	...	0.05
...	...	...	...	...	...
20150102	11:59:59:43	12345	1.6	...	0.06
...	...	...	...	...	...
20151231	00:00:01:43	12345	1.3	...	0.06
20151231	00:00:02:43	12345	1.3	...	0.05
...	...	...	...	...	...
20151231	11:59:59:43	12345	1.4	...	0.05

# Which problems go away with Kudu?





# Changing hardware landscape

- **Spinning disk -> solid state storage**
  - **NAND flash:** Up to 450k read 250k write iops, about 2GB/sec read and 1.5GB/sec write throughput, at a price of less than \$3/GB and dropping
  - **3D XPoint memory** (1000x faster than NAND, cheaper than RAM)
- **RAM** is cheaper and more abundant:
  - 64->128->256GB over last few years
- **Takeaway:** The **next bottleneck is CPU**, and current storage systems weren't designed with CPU efficiency in mind.

# Kudu CPU efficiency

- C++ implementation designed for modern CPUs
  - Use SSE instructions, optimize for deep CPU cache hierarchy
  - Embed Impala's LLVM-based query fragment execution for SQL access
  - Avoid GC and other Java pitfalls
- File formats designed for efficient deserialization
  - Type-specific compression much faster than generic (LZO)
  - Trade-offs: 10% space increase is worth it, if 100% faster to decode

# Kudu Basic Design

- Basic Construct: Tables
  - Tables broken down into a storage mechanism named Tablets (roughly equivalent to regions or partitions)
  - Columnar storage
  - Flexible data partitioning
- Typed storage
  - Enables compression
- Next generation consistency and replication
  - Multi-Version Concurrency Control (MVCC)
  - Raft Consensus to replicate *operations*
  - Architecture supports geographically disparate, active/active systems
    - Not in the initial implementation

# Kudu: Scalable and fast tabular storage

---

- **Scalable**

- Tested up to 275 nodes (~3PB cluster)
- Designed to scale to **1000s of nodes, tens of PBs**

- **Fast**

- **Millions** of read/write operations per second across cluster
- **Multiple GB/second** read throughput per node

- **Tabular**

- Store tables like a normal database (can support SQL, Spark, etc)
- Individual record-level access to **100+ billion row tables** (Java/C++/Python APIs)




# TPC-H (analytics benchmark)

- 75 server cluster
  - 12 (spinning) disks each, enough RAM to fit dataset
  - TPC-H Scale Factor 100 (100GB)

- Example query:

```
SELECT n_name, sum(l_extendedprice * (1 - l_discount)) as revenue
FROM customer, orders, lineitem, supplier, nation, region
WHERE c_custkey = o_custkey AND l_orderkey = o_orderkey
      AND l_suppkey = s_suppkey AND c_nationkey = s_nationkey
      AND s_nationkey = n_nationkey
      AND n_regionkey = r_regionkey
      AND r_name = 'ASIA'
      AND o_orderdate >= date '1994-01-01'
      AND o_orderdate < '1995-01-01'
GROUP BY n_name ORDER BY revenue desc;
```


## TPC-H SF 100 @75 nodes



- Kudu outperforms Parquet by 31% (geometric mean) for RAM-resident data


# Versus other NoSQL storage

- **Apache Phoenix: OLTP SQL engine built on HBase**
- 10 node cluster (9 worker, 1 master)
- TPC-H LINEITEM table only (6B rows)



# What about NoSQL-style random access? (YCSB)

- **YCSB** 0.5.0-snapshot
- 10 node cluster  
(9 worker, 1 master)
- 100M row data set
- 10M operations each workload





- Current trends
- Medical devices
- Infrastructure
- **Outcome**

# Parkinson Disease



- 60,000 Americans are diagnosed with Parkinson's every year
- No cure, several therapies with varying tradeoffs
- No objective measurement on therapy impact

# Therapies

- Drugs are minimally invasive, but can make symptoms worse and cause seizures
- Surgery to remove working portions of the brain
- Implantable medical devices are invasive and carry their own set of risks



# Solution





# Data driven disease management

- Device continually uploads tremor data
- Useful for multiple
  - Clinician
  - Researcher
  - Pharmaceutical & Medical Device
- Intel and Michael J. Fox Foundation
  - <http://goo.gl/Eccu5x>
  - <http://goo.gl/8nVPuU>

# Summary

- Mobile technology
- Medical devices
- Combined with open source big data technologies
  - Hadoop
  - Spark
  - Kafka
  - Kudu
- Results in healthier population and better care

# Brock Noland

brock@phdata.io  
@phdatainc  
phdata.io

